

## Project 3.A

He Zhu, hz6627, UT Austin

This project can be divided into 2 parts, one part is to implement null pointer dereference and another part is to implement the `mprotect()` and `munprotect()`. In this report, I presented several key steps to finish it.

To implement the null pointer, I tried to leave the first page table blank and boot the xv6 from the second page 0x1000. To implement this, the first thing I do is do some modification on Makefile:

```
_%.o: %.o $(ULIB)
$(LD) $(LDFLAGS) -N -e main -Ttext 0 -o $$ $^
```

should be modified to

Here I changed entry point from 0 to 0x1000 to make the first page invalid.

```
_%.o: %.o $(ULIB)
$(LD) $(LDFLAGS) -N -e main -Ttext 0x1000 -o $$ $^
```

After updating the Makefile, I looked into `exec()` and `fork()`. I firstly trace back to file `exec.c` to look into `exec()`.

```
// Load program into memory.
sz = 0;
for(i=0, off=elf.phoff; i<elf.phnum; i++, off+=sizeof(ph)){
    if(readi(ip, (char*)&ph, off, sizeof(ph)) != sizeof(ph))
        goto bad;
    if(ph.type != ELF_PROG_LOAD)
        continue;
    if(ph.memsz < ph.filesz)
        goto bad;
    if((sz = allocuvm(pgdir, sz, ph.vaddr + ph.memsz)) == 0)
        goto bad;
    if(loaduvm(pgdir, (char*)ph.vaddr, ip, ph.off, ph.filesz) < 0)
        goto bad;
}
```

The `sz` is assigned with `PGSIZE`, it is defined in the `mmu.h`, the header file contains the x86 memory management unit definitions, and `PGSIZE = 4096`, the same as 0x1000, and it means that when we start our program, the beginning memory point of the program becomes second page.

I also made sure when we worked on the virtual memory and physical memory, beginning page is well defined, so `allocuvm()`, `loaduvm()`, `switchuvm()`, `setupkvm()` are also carefully looked into.

```

// Load program into memory.
sz = PGSIZE;
for(i=0, off=elf.phoff; i<elf.phnum; i++, off+=sizeof(ph)){
    if(readi(ip, (char*)&ph, off, sizeof(ph)) != sizeof(ph))
        goto bad;
    if(ph.type != ELF_PROG_LOAD)
        continue;
    if(ph.memsz < ph.filesz)
        goto bad;
    if((sz = allocuvm(pgdir, sz, ph.vaddr + ph.memsz)) == 0)
        goto bad;
    if(loaduvm(pgdir, (char*)ph.vaddr, ip, ph.off, ph.filesz) < 0)
        goto bad;
}

```

Fortunately, they are not relevant with beginning page of the program, so we do not need to change anything. Then according to the project instruction, I looked into the `fork()`, which is defined in file `proc.c`, and here is one function `copyuvm()` used, I needed to do some modification on it, and `copyuvm()` is also defined in `vm.c`:

```

.....
for (i = 0; i < sz; i += PGSIZE)
{
    if ((pte = walkpgdir(pgdir, (void *)i, 0)) == 0)
        panic("copyuvm: pte should exist");
    if (!(*pte & PTE_P))
        panic("copyuvm: page not present");
    .....
}

```

should be changed to

```

.....
for (i = PGSIZE; i < sz; i += PGSIZE)
{
    if ((pte = walkpgdir(pgdir, (void *)i, 0)) == 0)
        panic("copyuvm: pte should exist");
    if (!(*pte & PTE_P))
        panic("copyuvm: page not present");
    .....
}

```

This means when we walk through the page directory table, we start from virtual address `0x1000`, the second memory page, thus ensuring when we invoked `fork()` to create a child process, the virtual memory is accurately copied. (starting from second page and end to the process size limit in bytes). Finally, I googled and found that I should do some optimization on functions `fetchstr()` and `fetchint()` of `syscall.c` file, but I did not show them here.

The second part of the project 3A is only to add kernel level functions and relevant user functions. Because I have previous done a project of adding a system call, here I just presented the core code segments I did:

sysproc.c

```
void _mprotect(struct proc *p, void *addr, int len) {  
    uint vpn, ad = (uint)addr;  
    int size = ad + len - 1;  
    for (vpn = ad; vpn < size; vpn += PGSIZE) {  
        pte_t *pte;  
        pde_t *pde = p->pgdir;  
        if ((pte = walkpgdir(pde, (void*)vpn, 0)) == 0) {  
        } else {  
            if ((*pte)&PTE_W && (*pte&PTE_U) == 0) {  
                *pte = (*pte)&(~PTE_W);  
                lcr3(v2p(proc->pgdir));  
            }  
        }  
    }  
    cprintf("\n");  
}
```

```
void _munprotect(struct proc *p, void *addr, int len) {  
    uint vpn, ad = (uint)addr;  
    int size = ad + (len * PGSIZE) + 1;  
    for (vpn = ad; vpn < size; vpn += PGSIZE) {  
        pte_t *pte;  
        pde_t *pde = p->pgdir;  
        if ((pte = walkpgdir(pde, (void*)vpn, 0)) == 0) {  
        } else {  
            if (!(*pte&PTE_W) && (*pte&PTE_U) == 0) {  
                *pte = *pte|PTE_W;  
                lcr3(v2p(proc->pgdir));  
            }  
        }  
    }  
    cprintf("\n");  
}
```

```

int kern_mprotect(void *addr, int len) {
    int rv = -1;
    if ((int)addr%PGSIZE != 0 || (int)addr > proc->sz || (int)addr <= 0) {
        return rv;
    }
    if (len <= 0 || ((int)addr + (len * PGSIZE)) > proc->sz) {
        return rv;
    }
    _mprotect(proc, addr, len);
    rv = 0;
    return rv;
}

int kern_munprotect(void *addr, int len) {
    int rv = -1;
    if ((int)addr%PGSIZE != 0 || (int)addr > proc->sz || (int)addr <= 0) {
        return rv;
    }
    if (len <= 0 || ((int)addr + (len*PGSIZE)) > proc->sz) {
        return rv;
    }
    _munprotect(proc, addr, len);
    rv = 0;
    return rv;
}

```

### Project 3.B

In this project, I implemented clone () and join () as system calls, the core codes are attached below:

```

int sys_clone(void) {
    {
        void *arg1, *arg2, *stk;
        void(*fnc) (void *, void *);
        if(argptr(0, (void *)&fnc, 0) < 0)
            return -1;
        if(argptr(1, (void *)&arg1, sizeof(void*)) < 0)
            return -1;
        if(argptr(2, (void *)&arg2, sizeof(void*)) < 0)
            return -1;
        if(argptr(3, (void *)&stk, PGSIZE) < 0)
            return -1;
        if((uint)stk % PGSIZE)
            return -1;
        return clone(fnc, arg1, arg2, stk);
    }
}

```

```

int sys_join(void)
{
    void **stackPointer = 0;

    if(argptr(0, (void*)&stackPointer, sizeof(void *)) < 0)
        return -1;

    return join(stackPointer);
}

```

The user-level calls are defined in file *proc.c*

```

int join(void **stkp)
{
    struct proc* proc = myproc();
    struct proc *p;
    int thexist, pid;
    acquire(&ptable.lock);
    while(1){
        thexist = 0;
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(proc->pgdir != p->pgdir || proc->pid == p->pid)
                continue;
            thexist = 1;
            if(p->state == ZOMBIE){
                pid = p->pid;
                kfree(p->kstack);
                p->pid = 0;
                p->parent = 0;
                p->killed = 0;
                p->name[0] = 0;
                p->kstack = 0;
                p->state = UNUSED;
                *stkp = p->stack;
                release(&ptable.lock);
                return pid;
            }
        }
        if(!(thexist && proc->killed)){
            release(&ptable.lock);
            return -1;
        }
        sleep(proc, 0);
    }
}

```

```

int clone(void(*func) (void *, void *), void*
arg1, void* arg2, void* stack)
{
    struct proc *proc = myproc();
    int pid, i;
    struct proc *np;
    char *stackaddr;
    if(!(np = allocproc()))
        return -1;
    np->pgdir = proc->pgdir;
    np->sz = proc->sz;
    np->parent = proc;
    *np->tf = *proc->tf;
    np->tf->eax = 0;
    np->tf->ebp = 0;
    np->tf->eip = (int) func;
    np->stack = stack;
    np->tf->esp = (int) (stack) + PGSIZE - 12;
    stackaddr = uva2ka(np->pgdir, (char*)stack);
    *(void **) (stackaddr + PGSIZE - 12) =
(void*) 0xffffffff;
    *(void **) (stackaddr + PGSIZE - 8) = (void*) arg1;
    *(void **) (stackaddr + PGSIZE - 4) = (void*) arg2;

    for(i = 0; i < NOFILE; i++)
        if(proc->ofile[i])
            np->ofile[i] = filedup(proc->ofile[i]);
    np->cwd = idup(proc->cwd);
    pid = np->pid;
    np->state = RUNNABLE;
    safestrcpy(np->name, proc->name,
sizeof(proc->name));
    return pid;
}

```

I also implemented several frequently used user-level functions:

```
void lock_init(lock_t *lock) {
    lock->ticket = 0;
    lock->turn = 0;
}

void lock_acquire(lock_t *lock) {
    int myturn = fadd(&lock->ticket, 1);
    while (lock->turn != myturn);
}

void lock_release(lock_t *lock) {
    fadd(&lock->turn, 1);
}
```

```
int
thread_join()
{
    void *stk;
    int cpid;
    //Need to free stack
    cpid = join(&stk);
    if (cpid != -1) {
        for(int i=0; i<ARRSZ; i++) {
            if(val_array[i] == cpid) {
                free(userstkaddr[i]);
                userstkaddr[i] = 0;
                val_array[i] = -1;
            }
        }
    }
    return cpid;
}
```

```
#define ARRSZ 64
int val_array[ARRSZ];
void* userstkaddr[ARRSZ];
int thread_create(void
(*start_routine)(void *, void *),
void *arg1, void *arg2)
{
    void* _stkaddr, *_stkpassed;
    _stkaddr = malloc(2*PGSIZE);
    if(_stkaddr == 0) {
        return -1;
    }
    //align the page
    int extspace =
((int) (_stkaddr))%PGSIZE;
    _stkpassed = (_stkaddr) + PGSIZE -
extspace;
    int cpid = clone(start_routine,
arg1, arg2, _stkpassed);
    if (cpid != -1) {
        for(int i=0; i<ARRSZ; i++) {
            if(userstkaddr[i] == 0) {
                userstkaddr[i] = _stkaddr;
                val_array[i] = cpid;
                break;
            }
        }
    }
    else {
        free(_stkaddr);
    }
    return cpid;
}
```